

In [23]:

```
import torchvision
import torch
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
```

In [3]:

```
# Data download and preprocessing

DOWNLOAD_MNIST = True # If already download , set as False
train_data = torchvision.datasets.MNIST(
    root = './mnist/',
    train = True , # this is training data
    # transform = torchvision.transforms.ToTensor(),
    download = DOWNLOAD_MNIST,
)
test_data = torchvision.datasets.MNIST(root = './mnist/', train = False)

# change the features to numpy
X_train = train_data.train_data.numpy()
X_test = test_data.test_data.numpy()

# change the labels to numpy
Y_train = train_data.train_labels.numpy()
Y_test = test_data.test_labels.numpy()
```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./mnist/MNIST/raw/train-images-idx3-ubyte.gz

Extracting ./mnist/MNIST/raw/train-images-idx3-ubyte.gz to ./mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./mnist/MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./mnist/MNIST/raw/train-labels-idx1-ubyte.gz to ./mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./mnist/MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./mnist/MNIST/raw/t10k-images-idx3-ubyte.gz to ./mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./mnist/MNIST/raw

/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:75: UserWarning: train_data has been renamed data
  warnings.warn("train_data has been renamed data")
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:80: UserWarning: test_data has been renamed data
  warnings.warn("test_data has been renamed data")
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:65: UserWarning: train_labels has been renamed targets
  warnings.warn("train_labels has been renamed targets")
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:70: UserWarning: test_labels has been renamed targets
  warnings.warn("test_labels has been renamed targets")

```

3(a)

In [24]:

```

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow.keras as keras
from keras.models import Sequential
from keras.layers import Dense
from keras.metrics import SparseCategoricalCrossentropy
from sklearn import preprocessing

```

In [5]:

```

mms = preprocessing.MinMaxScaler()
scaled_X_train = mms.fit_transform(X_train.reshape(60000,28*28))
X_train_new = scaled_X_train.reshape(60000, 784)

```


In []:

```

np.random.seed(1)
sgd=tf.keras.optimizers.SGD(learning_rate=0.1)
model= Sequential()
model.add(Dense(100,activation='relu',input_shape=(784,)))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=sgd,loss=tf.keras.losses.SparseCategoricalCrossentropy)
history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch_s

plt.plot(history.history['loss'], color='b', label="Training loss")
plt.plot(history.history['val_loss'], color='r', label="validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'],loc='best')
plt.show()

```




In []:

```

np.random.seed(2)
sgd=tf.keras.optimizers.SGD(learning_rate=0.1)
model= Sequential()
model.add(Dense(100,activation='relu',input_shape=(784,)))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=sgd,loss=tf.keras.losses.SparseCategoricalCrossentropy)
history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch_s

plt.plot(history.history['loss'], color='b', label="Training loss")
plt.plot(history.history['val_loss'], color='r', label="validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'],loc='best')
plt.show()

```




In []:

```

np.random.seed(3)
sgd=tf.keras.optimizers.SGD(learning_rate=0.1)
model= Sequential()
model.add(Dense(100,activation='relu',input_shape=(784,)))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=sgd,loss=tf.keras.losses.SparseCategoricalCrossentropy)
history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch_s

plt.plot(history.history['loss'], color='b', label="Training loss")
plt.plot(history.history['val_loss'], color='r', label="validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'],loc='best')
plt.show()

```




In []:

```

np.random.seed(4)
sgd=tf.keras.optimizers.SGD(learning_rate=0.1)
model= Sequential()
model.add(Dense(100,activation='relu',input_shape=(784,)))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=sgd,loss=tf.keras.losses.SparseCategoricalCrossentropy)
history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch_s

plt.plot(history.history['loss'], color='b', label="Training loss")
plt.plot(history.history['val_loss'], color='r', label="validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'],loc='best')
plt.show()

```




In []:

```

np.random.seed(5)
sgd=tf.keras.optimizers.SGD(learning_rate=0.1)
model= Sequential()
model.add(Dense(100,activation='relu',input_shape=(784,)))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=sgd,loss=tf.keras.losses.SparseCategoricalCrossentropy)
history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch_s

plt.plot(history.history['loss'], color='b', label="Training loss")
plt.plot(history.history['val_loss'], color='r', label="validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'],loc='best')
plt.show()

```





The validation error is higher than the training error. With the decrease of the training error, the validation error at first decrease and then a little increase, which suggests it's overfitting.

3(b)

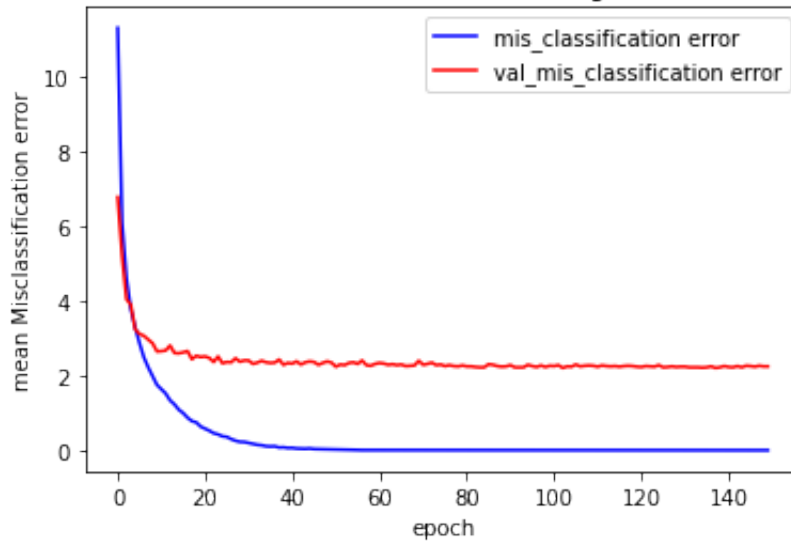
In []:

```
for seed in range(1,6):
    model = Sequential()
    model.add(Dense(100,activation='relu',input_shape=(784,)))
    model.add(Dense(10,activation='softmax'))
    np.random.seed(seed)
    model.compile(optimizer=sgd,loss=keras.losses.SparseCategoricalCrossentropy)
    history=model.fit(X_train_new,Y_train,epochs=150,validation_split=0.2,batch
    x = np.repeat(1,150)
    plt.plot((x-history.history['accuracy'])*100, color='b', label="mis_classif
    plt.plot((x-history.history['val_accuracy'])*100, color='r', label="val_mis
    plt.xlabel('epoch')
    plt.ylabel('mean Misclassification error')
    plt.title('mean Misclassification error using seed=%d'%(seed))
    plt.legend(['mis_classification error','val_mis_classification error'],loc=
    plt.show()
```

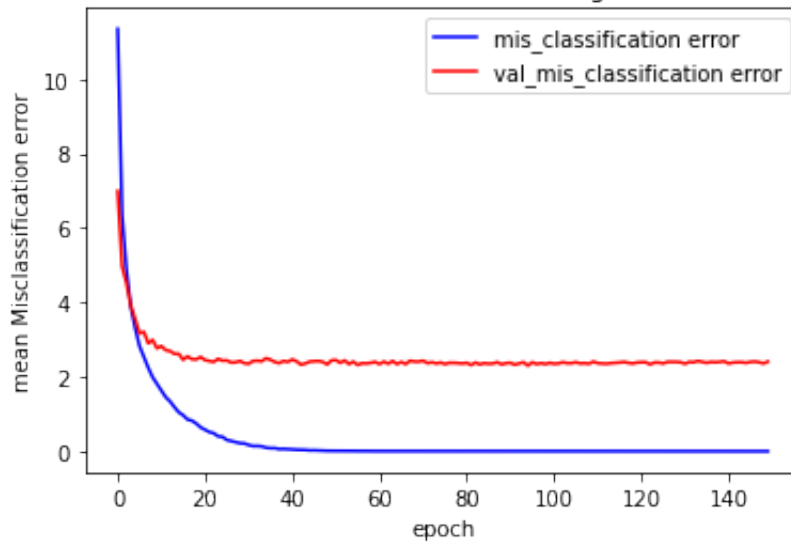






mean Misclassification error using seed=2



mean Misclassification error using seed=3

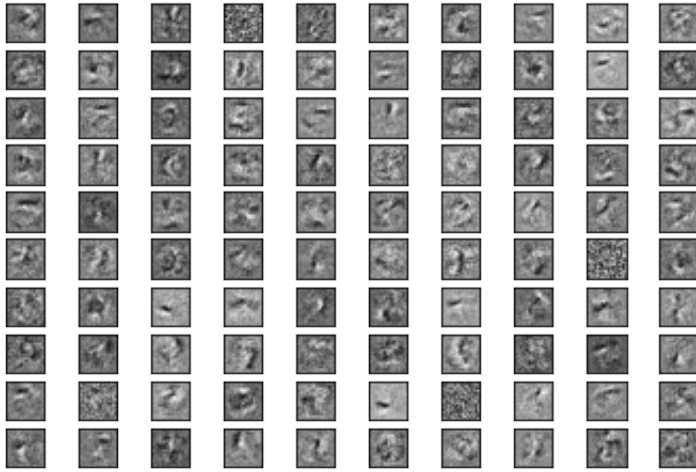




The validation error is still higher than the training error, but it is more certain compared to the behavior of the cross-entropy error function.

3(c)

```
In [ ]: W = model.layers[0].get_weights()[0].reshape(28,28,100)
for i in range(W.shape[2]):
    plt.subplot(10, 10, i + 1)
    plt.imshow(W[:, :, i], cmap = 'gray')
    plt.xticks([])
    plt.yticks([])
plt.show()
```

The plot shows that it is uncorrelated and not too noisy, as well as given some curvy shapes.

3(d)

In [6]:

```
scaled_X_test = mms.fit_transform(X_test.reshape(10000, 28*28))  
X_test_new = scaled_X_test.reshape(10000, 784)
```

In [7]:

```

seed=3

for lr in [0.01, 0.1, 0.2, 0.5]:
    for momentum in [0, 0.5, 0.9]:
        sgd=keras.optimizers.SGD(learning_rate = lr, momentum = momentum)
        np.random.seed(seed)

        #Create single layer ANN
        model = Sequential()
        #Adding Hidden layer
        model.add(Dense(units=100, activation='relu', input_shape=(784,)))
        #Adding Output layer
        model.add(Dense(units=10, activation='softmax'))
        #Compiling the model
        model.compile(optimizer=sgd, loss=keras.losses.SparseCategoricalCrossentropy,
                      metrics=['accuracy'])
        history = model.fit(X_train_new, Y_train, epochs=150,
                           validation_data=(X_test_new, Y_test), batch_size=64, verbose=0)


        #Plot
        f, ax = plt.subplots(1, 2, figsize=(15, 5))
        plt.subplots_adjust(left=0.1,
                           bottom=0.1,
                           right=0.9,
                           top=0.9,
                           wspace=0.5,
                           hspace=0.5)



        ax[0].plot(history.history['loss'], color='blue', label="Training loss")
        ax[0].plot(history.history['val_loss'], color='red', label="validation loss")
        ax[0].set_xlabel('epoch')
        ax[0].set_ylabel('average CrossEntropy')
        ax[0].set_title('Cross Entropy')
        ax[0].legend(['loss', 'val_loss'], fontsize = 10)

        x = np.repeat(1, 150)
        ax[1].plot((x - history.history['accuracy'])*100, color = 'blue', label = 'mis_classification_error')
        ax[1].plot((x - history.history['val_accuracy'])*100, color = 'red', label = 'val_mis_classification_error')
        ax[1].set_xlabel('epoch')
        ax[1].set_ylabel('mean Misclassification error')
        ax[1].set_title('mean Misclassification error')
        ax[1].legend(['mis_classification_error', 'val_mis_classification_error'])

        f.suptitle('Learning rate=%2f & momentum=%2f'%(lr, momentum))
        plt.show()
        print("The misclassification error for test set is %3f"%((x - history.history['val_accuracy'])*100))

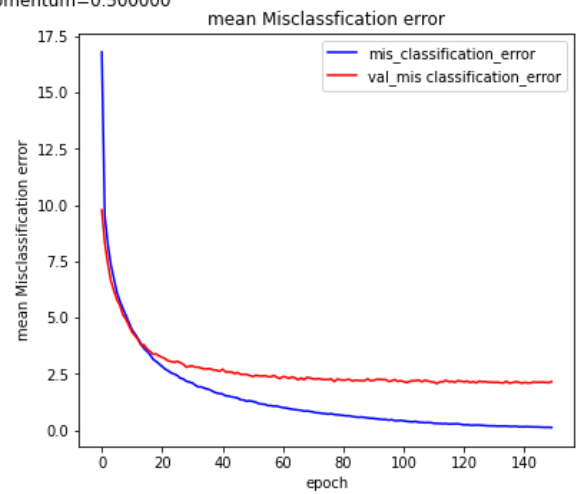

```






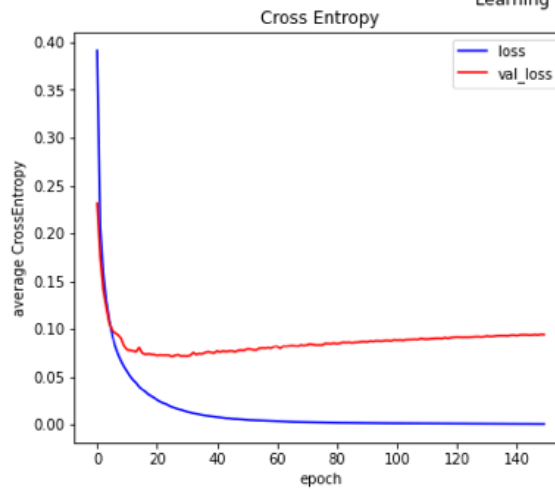
The misclassification error for test set is 6.230003%.

Learning rate=0.010000 & momentum=0.500000

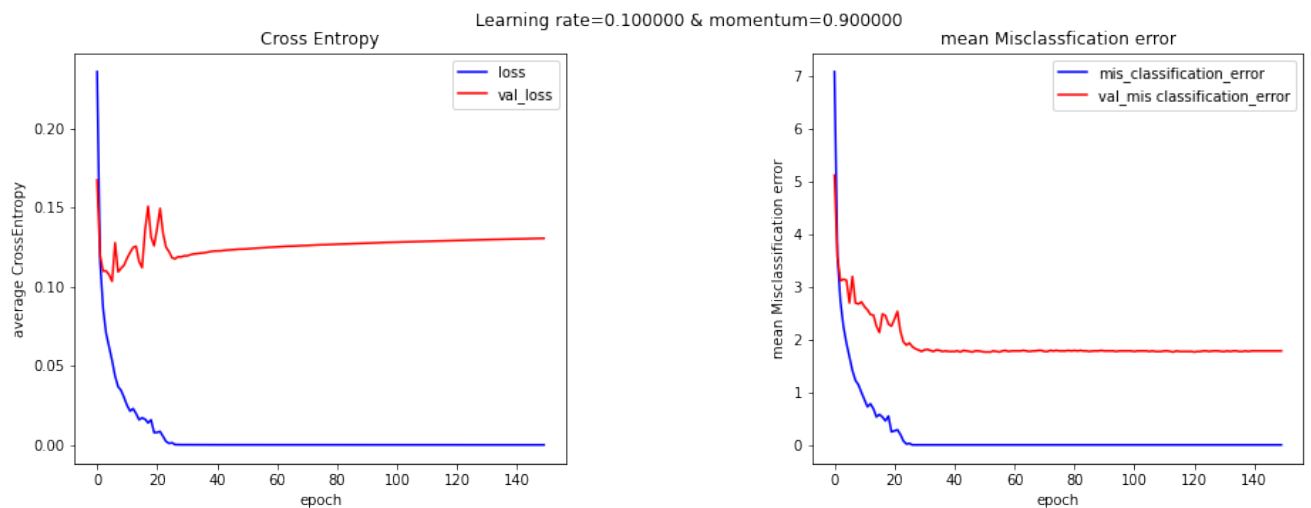

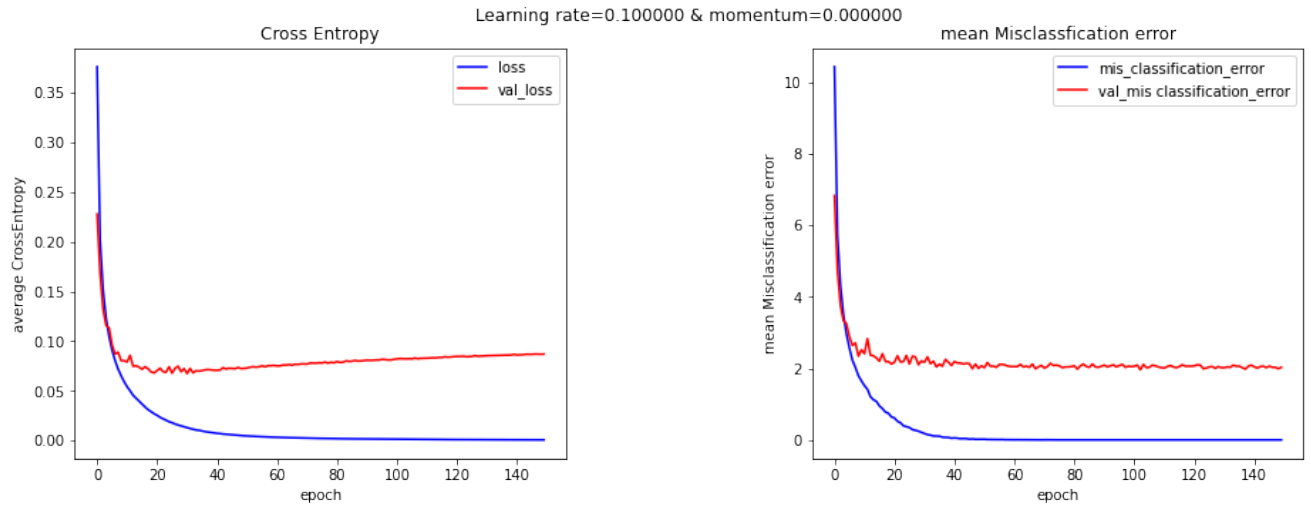





The misclassification error for test set is 4.640001%.


Learning rate=0.010000 & momentum=0.900000




The misclassification error for test set is 2.450001%.








The misclassification error for test set is 2.319998%.



The misclassification error for test set is 2.660000%.



The misclassification error for test set is 72.229999%.

From plots, if we fix the learning rate, we can see that when momentum is increasing, the error rate in most of cases is likely decreasing. If we fix the momentum, we can see that when learning rate is increasing, the error rate becomes more variation. The best value of these parameters seems like learning rate = 0.1 and momentum = 0.5,

4(a)

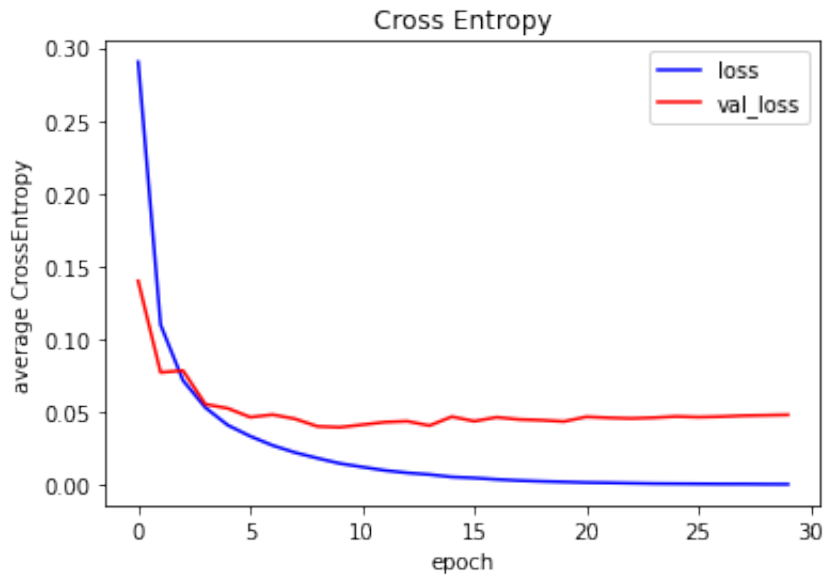
```
In [8]: #Since 2-D space, reshape the data
X_train_new1 = X_train_new.reshape(X_train.shape[0], 28, 28, 1)
X_test_new1 = X_test_new.reshape(X_test.shape[0], 28, 28, 1)
```

```
In [9]: from keras.layers import Conv2D,MaxPooling2D,Flatten
```

```
In [ ]: np.random.seed(3)
sgd = keras.optimizers.SGD(learning_rate = 0.1)
#Create single layer NN
model = Sequential()
#Adding Hidden layer
model.add(Conv2D(32, (3, 3), padding = 'same', activation='relu', input_shape=
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(units = 100, activation = 'relu'))
model.add(Dense(units = 10, activation = 'softmax'))
#Compiling the model
model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossentr
history = model.fit(X_train_new1, Y_train, epochs = 30, validation_data = (X_

#Plot
plt.plot(history.history['loss'], color = 'blue', label = "Training loss")
plt.plot(history.history['val_loss'], color = 'red', label = "validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'], fontsize = 10)

plt.show()
```



4(b)

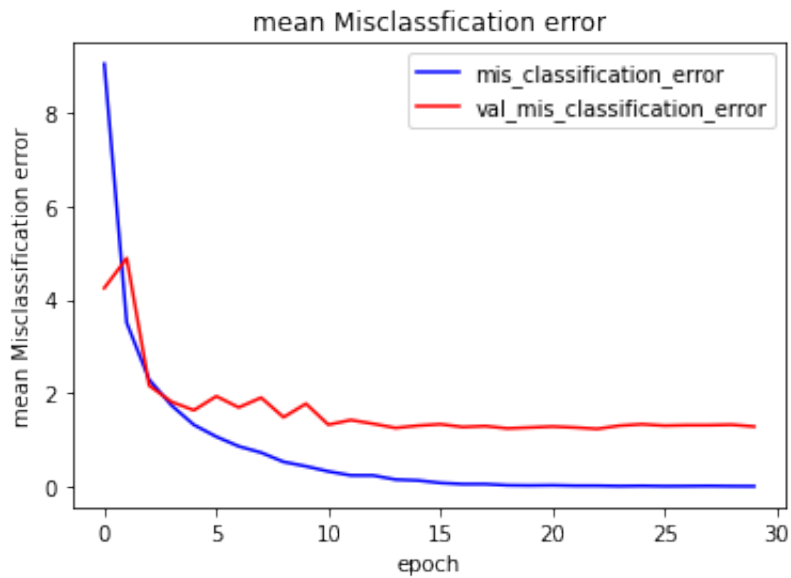
In []:

```

np.random.seed(3)
sgd = keras.optimizers.SGD(learning_rate = 0.1)
#Create single layer NN
model = Sequential()
#Adding Hidden layer
model.add(Conv2D(32, (3, 3), padding = 'same', activation='relu', input_shape=
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(units = 100, activation = 'relu'))
model.add(Dense(units = 10, activation = 'softmax'))
#Compiling the model
model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossentr
history = model.fit(X_train_new1, Y_train, epochs = 30, validation_data = (X_

#Plot
x = np.repeat(1, 30)
plt.plot((x - history.history['accuracy'])*100, color = 'blue', label = "mis_
plt.plot((x - history.history['val_accuracy'])*100, color = 'red', label = "v
plt.xlabel('epoch')
plt.ylabel('mean Misclassification error')
plt.title('mean Misclassification error')
plt.legend(['mis_classification_error', 'val_mis_classification_error'], font
plt.show()


```

4(c)

In [13]:

```
W = model.layers[0].get_weights()[0].reshape(3, 3, -1)
for i in range(W.shape[2]):
    plt.subplot(4, 8, i + 1)
    plt.imshow(W[:, :, i], cmap = 'gray')
    plt.xticks([])
    plt.yticks([])
plt.show()
```



4(d)

In [10]:

```

seed = 3

for lr in [0.01, 0.1, 0.2, 0.5]:
    for momentum in [0, 0.5, 0.9]:
        sgd = keras.optimizers.SGD(learning_rate = lr, momentum = momentum)
        np.random.seed(seed)


        #Create single layer NN
        model = Sequential()
        #Adding Hidden layer
        model.add(Conv2D(32, (3, 3), padding = 'same', activation = 'relu', input_shape=(28, 28, 1)))
        model.add(MaxPooling2D((2, 2)))
        model.add(Flatten())
        model.add(Dense(units = 100, activation = 'relu'))
        model.add(Dense(units = 10, activation = 'softmax'))
        #Compiling the model
        model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossEntropy, metrics=['accuracy'])
        history = model.fit(X_train_new1, Y_train, epochs = 30, validation_data = (X_test_new1, Y_test))



        #Plot
        f, ax = plt.subplots(1, 2, figsize = (15, 5))
        plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.1)
        ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
        ax[0].plot(history.history['val_loss'], color = 'red', label = "validation loss")
        ax[0].set_xlabel('epoch')
        ax[0].set_ylabel('average CrossEntropy')
        ax[0].set_title('Cross Entropy')
        ax[0].legend(['loss', 'val_loss'], fontsize = 10)

        x = np.repeat(1, 30)
        ax[1].plot((x - history.history['accuracy'])*100, color = 'blue', label = "Training misclassification error")
        ax[1].plot((x - history.history['val_accuracy'])*100, color = 'red', label = "validation misclassification error")
        ax[1].set_xlabel('epoch')
        ax[1].set_ylabel('mean Misclassification error')
        ax[1].set_title('mean Misclassification error')
        ax[1].legend(['mis_classification_error', 'val_mis_classification_error'], fontsize = 10)

        f.suptitle('Learning rate=%2f & momentum=%2f'%(lr, momentum))
        plt.show()
        print("The misclassification error for test set is %3f"%((x - history.history['val_accuracy'])*100))



```







The misclassification error for test set is 3.259999%.

Learning rate=0.010000 & momentum=0.500000





The misclassification error for test set is 2.819997%.

Learning rate=0.010000 & momentum=0.900000




The misclassification error for test set is 1.470000%.


Learning rate=0.100000 & momentum=0.000000




The misclassification error for test set is 1.239997%.




The misclassification error for test set is 1.249999%.




The misclassification error for test set is 1.389998%.




The misclassification error for test set is 2.460003%.




The misclassification error for test set is 1.279998%.




The misclassification error for test set is 3.420001%.



The misclassification error for test set is 1.120001%.



The misclassification error for test set is 1.770002%.



The misclassification error for test set is 90.260000%.

From above, we can see that most of the set are stable when epoch > 20. When learning rate = 0.5 or momentum = 0.9, it is less stable than other parameters. And when rate = 0.5 and momentum = 0.9, it is fluctuated. The best value of the parameters of the single layer is learning rate = 0.1 and momentum = 0.5, with the misclassification error = 2.249998%, and for one layer of 2-D CNN, the misclassification error = 1.120001%, it clearly shows that one layer of 2-D CNN is outstanding. For parameters visualization, in the last row, most dark parts are distributed on sides. And it is like the overall trend.

5(a)

```
In [15]: from keras.layers import Dropout
```

In [15]:


```

np.random.seed(3)
sgd = keras.optimizers.SGD(learning_rate = 0.1)
#Create single layer NN
model = Sequential()
#Adding Hidden layer
model.add(Conv2D(32, (3, 3), padding = 'same', activation = 'relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.5))
model.add(Conv2D(64, (3, 3), padding = "same", activation = "relu"))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(units = 100, activation = 'relu'))
model.add(Dense(units = 10, activation = 'softmax'))
#Compiling the model
model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossentropy, metrics = ['accuracy'])
history = model.fit(X_train_new1, Y_train, epochs = 30, validation_data = (X_test, Y_test))

#Plot
plt.plot(history.history['loss'], color = 'blue', label = "Training loss")
plt.plot(history.history['val_loss'], color = 'red', label = "validation loss")
plt.xlabel('epoch')
plt.ylabel('average CrossEntropy')
plt.title('Cross Entropy')
plt.legend(['loss', 'val_loss'], fontsize = 10)

plt.show()

```




5 (b)

In [16]:

```

x = np.repeat(1, 30)
plt.plot((x - history.history['accuracy'])*100, color = 'blue', label = "mis_
plt.plot((x - history.history['val_accuracy'])*100, color = 'red', label = "v
plt.xlabel('epoch')
plt.ylabel('mean Misclassification error')
plt.title('mean Misclassification error')
plt.legend(['mis_classification_error', 'val_mis_classification_error'], font
plt.show()

```




5(c)

In [19]:

```

W = model.layers[0].get_weights()[0].reshape(3, 3, -1)
for i in range(W.shape[2]):
    plt.subplot(4, 8, i + 1)
    plt.imshow(W[:, :, i], cmap = 'gray')
    plt.xticks([])
    plt.yticks([])
plt.show()

```

5(d)

In []:

```

seed = 3

for lr in [0.01, 0.1, 0.2, 0.5]:
    for momentum in [0, 0.5, 0.9]:
        sgd = keras.optimizers.SGD(learning_rate = lr, momentum = momentum)
        np.random.seed(seed)


        #Create single layer NN
        model = Sequential()
        #Adding Hidden layer
        model.add(Conv2D(32, (3, 3), padding = 'same', activation = 'relu', input_shape=(28, 28, 1)))
        model.add(MaxPooling2D((2, 2)))
        model.add(Conv2D(64, (3, 3), padding = "same", activation = "relu"))
        model.add(MaxPooling2D((2, 2)))
        model.add(Dropout(0.5))
        model.add(Flatten())
        model.add(Dense(units = 100, activation = 'relu'))
        model.add(Dense(units = 10, activation = 'softmax'))
        #Compiling the model
        model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossEntropy, metrics=['accuracy'])
        history = model.fit(X_train_new1, Y_train, epochs = 30, validation_data = (X_test_new1, Y_test))



        #Plot
        f, ax = plt.subplots(1, 2, figsize = (15, 5))
        plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.1)
        ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
        ax[0].plot(history.history['val_loss'], color = 'red', label = "validation loss")
        ax[0].set_xlabel('epoch')
        ax[0].set_ylabel('average CrossEntropy')
        ax[0].set_title('Cross Entropy')
        ax[0].legend(['loss', 'val_loss'], fontsize = 10)

        x = np.repeat(1, 30)
        ax[1].plot((x - history.history['accuracy'])*100, color = 'blue', label = "Training misclassification error")
        ax[1].plot((x - history.history['val_accuracy'])*100, color = 'red', label = "validation misclassification error")
        ax[1].set_xlabel('epoch')
        ax[1].set_ylabel('mean Misclassification error')
        ax[1].set_title('mean Misclassification error')
        ax[1].legend(['mis_classification_error', 'val_mis classification_error'], fontsize = 10)

        f.suptitle('Learning rate=%2f & momentum=%2f'%(lr, momentum))
        plt.show()
        print("The misclassification error for test set is %3f"%((x - history.history['val_accuracy'])*100))



```







The misclassification error for test set is 1.889998%.

Learning rate=0.010000 & momentum=0.500000




The misclassification error for test set is 1.279998%.


Learning rate=0.010000 & momentum=0.900000




The misclassification error for test set is 0.720000%.






The misclassification error for test set is 0.800002%.




The misclassification error for test set is 0.849998%.




The misclassification error for test set is 2.029997%.





The misclassification error for test set is 0.980002%.



The misclassification error for test set is 89.720000%.

From the above, the error rate of most situations is smaller than one 2-D convolutional layer. The mean misclassification error is more stable when learning rate = 0.01. And when momentum equals to 0.9, no matter which learning rate, it is unstable. The one layer of 2-D convolutional has the best value of the parameters when learning rate = 0.5, momentum = 0.0 with lower misclassification error, which is 1.120001% while the two layers of 2-D convolutional has the best value of the parameters when learning rate = 0.2, momentum = 0.5 with lower misclassification error, which is 0.669998%. Since the misclassification error = 0.669998% is smaller than the performance of SVM with Gaussian Kernel, our deep learning architecture beats SVM with Gaussian Kernel. For parameters visualization, in the third row, upper is dark, which is not too noisy, and too correlated.

```
In [21]: # install pydrive to load data
!pip install -U -q PyDrive

from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials
import pandas as pd


auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)
```

```
In [25]: id = "1Qt0RPvkQW8JgM_2lAhT0fy4G-H67Jivq"
file = drive.CreateFile({'id':id})
file.GetContentFile('test.txt')
test = pd.read_csv('test.txt',header=None).to_numpy()
```

```
In [26]: id = "17qlY1tu2DwYZ-komyT957yqHxhuqIgk1"
file = drive.CreateFile({'id':id})
file.GetContentFile('train.txt')
train = pd.read_csv('train.txt',header=None).to_numpy()
```

```
In [27]: id = "18_OK5cP6bpLJ3yi_vhcFqrffwN2Ts_Bs"
file = drive.CreateFile({'id':id})
file.GetContentFile('val.txt')
val = pd.read_csv('val.txt',header=None).to_numpy()
```

```
In [27]: plt.figure(figsize=(10,10))
for i in range(25):
    s=train[i,:1568].reshape(28,56)
    plt.subplot(5,5,i+1)
    plt.grid(False)
    plt.imshow(s, cmap=plt.cm.binary)
plt.show()
```



```
In [28]: print(train[0:25,-1].reshape(5,5))
```

```
[[ 5.  5.  4. 16. 12.]
 [15.  9.  8.  1.  5.]
 [13. 10. 12. 11.  5.]
 [11.  8. 14.  5.  3.]
 [ 2. 12. 14. 13.  9.]]
```

The pixels were scanned out in row-major and the relationship between the 2 digits and the last coordinate of each line is that the sum of the 2 digits is equal to the value of the last coordinate of each line

7

first:

In [28]:

```
X_train = train[:, :1568].reshape((-1, 28, 56, 1))
Y_train = train[:, -1]
X_val = val[:, :1568].reshape((-1, 28, 56, 1))
Y_val = val[:, -1]
X_test = test[:, :1568].reshape((-1, 28, 56, 1))
Y_test = test[:, -1]
```

In [19]:

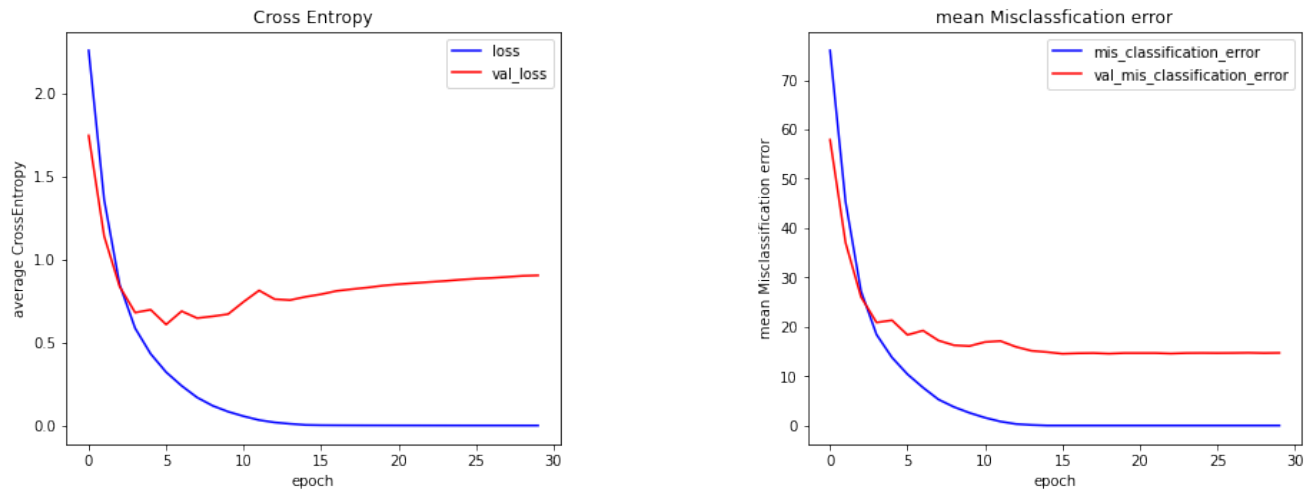
```
np.random.seed(3)
sgd = keras.optimizers.SGD(learning_rate = 0.1, momentum = 0.5)
#Create single layer NN
model = Sequential()
#Adding Hidden layers
model.add(Conv2D(32, (3, 3), padding = 'same', activation='relu', input_shape=(28, 56, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(units = 64, activation = 'relu'))
model.add(Dense(units = 19))
#Compiling the model
model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossentropy, metrics=['accuracy'])
history = model.fit(X_train, Y_train, epochs = 30, validation_data = (X_val, Y_val))

#Plot
f, ax = plt.subplots(1, 2, figsize = (15, 5))
plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.5)
ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
ax[0].plot(history.history['val_loss'], color = 'red', label = "validation loss")
ax[0].set_xlabel('epoch')
ax[0].set_ylabel('average CrossEntropy')
ax[0].set_title('Cross Entropy')
ax[0].legend(['loss', 'val_loss'], fontsize = 10)

x = np.repeat(1, 30)
ax[1].plot((x - history.history['accuracy'])*100, color = 'blue', label = "misclassification error")
ax[1].plot((x - history.history['val_accuracy'])*100, color = 'red', label = "validation misclassification error")
ax[1].set_xlabel('epoch')
ax[1].set_ylabel('mean Misclassification error')
ax[1].set_title('mean Misclassification error')
ax[1].legend(['mis_classification_error', 'val_mis_classification_error'], fontsize = 10)


plt.show()

pred = np.argmax(model.predict(X_test), axis=1)
err = np.mean(pred != Y_test)*100
print("The misclassification error for validation set and test set are %3f"%(history.history['val_loss'])
      'and %3f'%(err)+'%.')
```



The misclassification error for validation set and test set are 16.100001% and 13.180000%.

```
In [30]: W = model.layers[0].get_weights()[0].reshape(3, 3, -1)
for i in range(W.shape[2]):
    plt.subplot(4, 8, i + 1)
    plt.imshow(W[:, :, i], cmap = 'gray')
    plt.xticks([])
    plt.yticks([])
plt.show()
```



In [21]:

```

seed = 3

for lr in [0.01, 0.1, 0.2, 0.5]:
    for momentum in [0, 0.5, 0.9]:
        sgd = keras.optimizers.SGD(learning_rate = lr, momentum = momentum)
        np.random.seed(seed)


        #Create single layer NN
        model = Sequential()
        #Adding Hidden layer
        model.add(Conv2D(32, (3, 3), padding = 'same', activation = 'relu', input_shape=(X_train.shape[1], X_train.shape[2], X_train.shape[3])))
        model.add(MaxPooling2D((2, 2)))
        model.add(Flatten())
        model.add(Dense(units = 64, activation = 'relu'))
        model.add(Dense(units = 19))
        #Compiling the model
        model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossEntropy, metrics = ['accuracy'])
        history = model.fit(X_train, Y_train, epochs = 30, validation_data = (X_test, Y_test))



        #Plot
        f, ax = plt.subplots(1, 2, figsize = (15, 5))
        plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.1)
        ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
        ax[0].plot(history.history['val_loss'], color = 'orange', label = "validation loss")
        ax[0].set_xlabel('epoch')
        ax[0].set_ylabel('average CrossEntropy')
        ax[0].set_title('Cross Entropy')
        ax[0].legend(['loss', 'val_loss'], fontsize = 10)

        x = np.repeat(1, 30)
        ax[1].plot(x - history.history['accuracy']*100, color = 'blue', label = "Training accuracy")
        ax[1].plot(x - history.history['val_accuracy']*100, color = 'orange', label = "validation accuracy")
        ax[1].set_xlabel('epoch')
        ax[1].set_ylabel('mean Misclassification error')
        ax[1].set_title('mean Misclassification error')
        ax[1].legend(['mis_classification_error', 'val_mis_classification_error'], fontsize = 10)


        f.suptitle('Learning rate=%2f & momentum=%2f'%(lr, momentum))
        plt.show()
        pred = np.argmax(model.predict(X_test), axis=1)
        err = np.mean(pred != Y_test)*100
        print("The misclassification error for validation set and test set are %3f and %3f"%(err)+'.')

```









The misclassification error for validation set and test set are 47.259998% and 29.560000%.









The misclassification error for validation set and test set are 37.260002% and 19.380000%.



The misclassification error for validation set and test set are 19.139999% and 14.820000%.







From above,we can see that some of mean misclassification error are very unstable, and are fluctuated. When momentum = 0.9, the fluctuation is much more serious. And some of learning rate and momentum sets face overfitting problem. Compared with the original data, the best value of the parameters is learning rate = 0.2, momentum = 0.5, with the misclassification error 14.539999% and 11.360000%. The misclassification error for validation set and test set are both higher than the original data. The largest misclassification error are 91.200000% and 90.980000%, where learning rate = 0.5, momentum = 0.9. For parameters visualization, most dark parts are distributed on sides. It is not too noisy, and too correlated.

second:

```
In [31]: from tensorflow.keras.layers import BatchNormalization
```


In [33]:

```

np.random.seed(3)
sgd = keras.optimizers.SGD(learning_rate = 0.1, momentum = 0.5)
#Create single layer NN
model = Sequential()
#Adding Hidden layers
model.add(Conv2D(32, (3, 3), padding = 'same', activation='relu', input_shape=(3, 3, 3)))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3), activation="relu"))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(units = 64, activation = 'relu'))
model.add(Dense(units = 19))
#Compiling the model
model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossentropy, metrics=['accuracy'])
history = model.fit(X_train, Y_train, epochs = 30, validation_data = (X_val, Y_val))

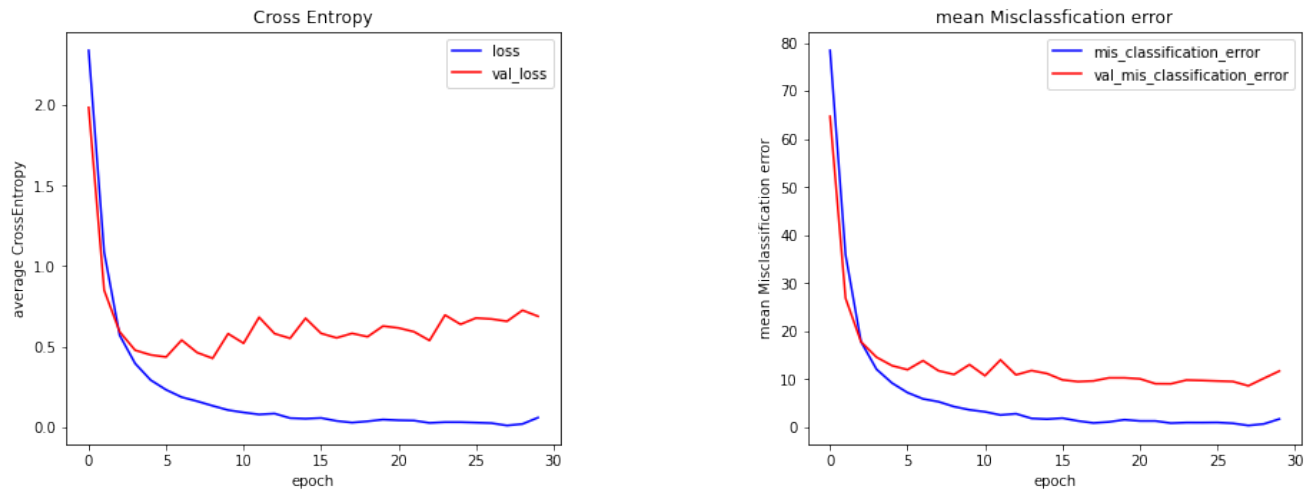
#Plot
f, ax = plt.subplots(1, 2, figsize = (15, 5))
plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.1)
ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
ax[0].plot(history.history['val_loss'], color = 'red', label = "validation loss")
ax[0].set_xlabel('epoch')
ax[0].set_ylabel('average CrossEntropy')
ax[0].set_title('Cross Entropy')
ax[0].legend(['loss', 'val_loss'], fontsize = 10)

x = np.repeat(1, 30)
ax[1].plot((x - history.history['accuracy'])*100, color = 'blue', label = "misclassification error")
ax[1].plot((x - history.history['val_accuracy'])*100, color = 'red', label = "validation misclassification error")
ax[1].set_xlabel('epoch')
ax[1].set_ylabel('mean Misclassification error')
ax[1].set_title('mean Misclassification error')
ax[1].legend(['mis_classification_error', 'val_mis_classification_error'], fontsize = 10)

plt.show()

pred = np.argmax(model.predict(X_test), axis=-1)
err = np.mean(pred != Y_test)*100
print("The misclassification error for validation set and test set are %3f"%(history.history['val_loss'])
      "and %3f"%(err)+"%.")


```



The misclassification error for validation set and test set are 13.000000% and 9.420000%.

In [34]:

```
W = model.layers[0].get_weights()[0].reshape(3, 3, -1)
for i in range(W.shape[2]):
    plt.subplot(4, 8, i + 1)
    plt.imshow(W[:, :, i], cmap = 'gray')
    plt.xticks([])
    plt.yticks([])
plt.show()
```



In [25]:

```

seed = 3

for lr in [0.01, 0.1, 0.2, 0.5]:
    for momentum in [0, 0.5, 0.9]:
        sgd = keras.optimizers.SGD(learning_rate = lr, momentum = momentum)
        np.random.seed(seed)

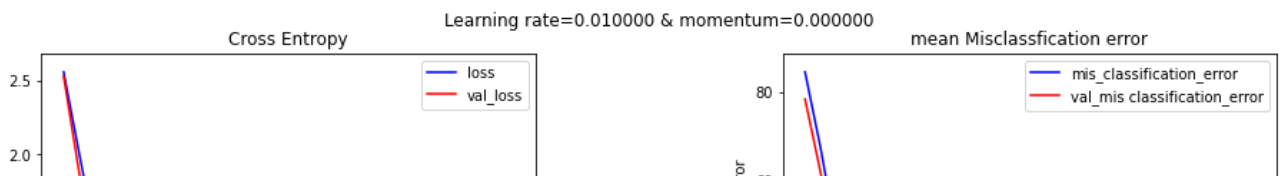
        #Create single layer NN
        model = Sequential()
        #Adding Hidden layer
        model.add(Conv2D(32, (3, 3), padding = 'same', activation = 'relu', input_shape=(X_train.shape[1], X_train.shape[2], X_train.shape[3])))
        model.add(MaxPooling2D((2, 2)))
        model.add(BatchNormalization())
        model.add(Conv2D(64, (3, 3), activation="relu"))
        model.add(MaxPooling2D((2, 2)))
        model.add(Flatten())
        model.add(Dense(units = 64, activation = 'relu'))
        model.add(Dense(units = 19))
        #Compiling the model
        model.compile(optimizer = sgd, loss = keras.losses.SparseCategoricalCrossEntropy, metrics=['accuracy'])
        history = model.fit(X_train, Y_train, epochs = 30, validation_data = (X_test, Y_test))



        #Plot
        f, ax = plt.subplots(1, 2, figsize = (15, 5))
        plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.5)
        ax[0].plot(history.history['loss'], color = 'blue', label = "Training loss")
        ax[0].plot(history.history['val_loss'], color = 'red', label = "validation loss")
        ax[0].set_xlabel('epoch')
        ax[0].set_ylabel('average CrossEntropy')
        ax[0].set_title('Cross Entropy')
        ax[0].legend(['loss', 'val_loss'], fontsize = 10)

        x = np.repeat(1, 30)
        ax[1].plot(x - history.history['accuracy']*100, color = 'blue', label = "Training accuracy")
        ax[1].plot(x - history.history['val_accuracy']*100, color = 'red', label = "validation accuracy")
        ax[1].set_xlabel('epoch')
        ax[1].set_ylabel('mean Misclassification error')
        ax[1].set_title('mean Misclassification error')
        ax[1].legend(['mis_classification_error', 'val_mis classification_error'], fontsize = 10)


        f.suptitle('Learning rate=%2f & momentum=%2f'%(lr, momentum))
        plt.show()
        pred = np.argmax(model.predict(X_test), axis=1)
        err = np.mean(pred != Y_test)*100
        print("The misclassification error for validation set and test set are %3f and %3f"%(err)+'.')

```









The misclassification error for validation set and test set are 16.140002% and 9.800000%.







The misclassification error for validation set and test set are 12.500000% and 9.240000%.




The misclassification error for validation set and test set are 11.379999% and 7.820000%.








The misclassification error for validation set and test set are 90.640000% and 89.300000%.



The misclassification error for validation set and test set are 90.640000% and 89.300000%.



The misclassification error for validation set and test set are 99.060000% and 99.120000%.

From above, we can see that most of mean misclassification error are very unstable and when learning rate grows, the unstable situation is more clear. It fluctuates between different epochs and some misclassification error lines for validation set and test set are even separate. Compared with the original data, the best value of the parameters is learning rate = 0.1, momentum = 0.0, with the misclassification error 9.799999% and 7.300000%. The misclassification error for validation set and test set are both much higher than the original data. For parameters visualization, most dark parts are distributed on sides. It is not too noisy, and too correlated.

why can't lower than 1%

Assuming that the neural network can perfectly calculate the sum of the number and label of two images, not affected by the stitching of two images. The lowest misclassification error 0.669998%, with learning rate = 0.2 and momentum = 0.0. Thus, the probability of two correct classification is $(1 - 0.669998\%)(1 - 0.669998\%) = 0.9866$. Then, even if the assumption holds, it still cannot obtain a test error lower than 1%.

Reference: <https://machinelearningmastery.com/how-to-visualize-filters-and-feature-maps-in-convolutional-neural-networks/>